

Five experiments: before and after GRPO

Final greedy accuracy on five generated evaluation batches, 500 questions per batch

● Base ◆ Trained



One saved training run per configuration. Descriptive means; no uncertainty for independent training repeats is available.