


Training stability on 6-digit addition

Full FT and LoRA configurations on the same task; other training hyperparameters also differ



Thin lines: raw logged values. Thick lines: trailing 20-step mean (starts at step 20).
Gradient norm is shown as logged; comparisons do not isolate the effect of LoRA.