



Reward recovery does not guarantee evaluation gains

3-digit addition | Full FT | training reward and final greedy evaluation are different measurements



Training: raw trajectory + trailing 20-step mean; reward includes correctness and format components (maximum 1.2).
Evaluation: five batches of 500 questions. No intermediate checkpoint evaluation is available.